

Cheetah provides orders of magnitude faster and differentiable beam dynamics simulations to speed up reinforcement learning and support applications like gradient-based optimisation.



Start using Cheetah today!

Scan for repository



Discord



```
pip install cheetah-accelerator

# Load initial beam distribution from ASTRA tracking
incoming = ParticleBeam.from_astra("beam_in.inl")

# Create a FODO lattice
segment = Segment(
    [
        Drift(length=torch.tensor(0.2)),
        Quadrupole(length=torch.tensor(0.2), name="Q1"),
        Drift(length=torch.tensor(0.4)),
        Quadrupole(length=torch.tensor(0.2), name="Q2"),
        Drift(length=torch.tensor(0.2)),
    ]
)

# Change the magnet strengths
segment.Q1.k1 = torch.tensor(10.0)
segment.Q2.k1 = torch.tensor(-9.0)

# Tracking through the segment
outgoing = segment.track(incoming)
```

Cheetah: A High-Speed Differentiable Beam Dynamics Simulator for Machine Learning Applications.

Jan Kaiser*, Chenran Xu, Annika Eichler, Andrea Santamaria Garcia

Motivation

- Machine learning has emerged as a powerful solution to the modern challenges in accelerator physics.
- Limited availability of beam time, the computational cost of simulations, and the high-dimensionality of optimisation problems pose significant challenges in generating the required data for training state-of-the-art machine learning models.

Beam dynamics simulations in PyTorch

- We introduce Cheetah, a beam dynamics simulation Python package written in PyTorch.
- Cheetah is up to 7 orders of magnitude faster than other simulation codes, making use of automatic and opt-in optimisations as well as PyTorch's built-in GPU acceleration.
- Being based on PyTorch, all models built in Cheetah support automatic differentiation, making Cheetah differentiable.

Example applications

- Two learning-based optimisation methods have found particular interest in the accelerator community:
 - A **reinforcement learning-trained optimisation** policy has been trained for 6 million interactions in feasible time using Cheetah and was successfully transferred zero-shot to the real ARES accelerator.
 - A **Gradient-based actuator tuning** on a Cheetah model can be used to tune the magnets at ARES to achieve desired transverse beam parameters.
 - A **Gradient-based system identification** can be used to build virtual diagnostics. At ARES for example, we can utilise Cheetah to identify quadrupole misalignments zero-shot from parasitic measurements.
 - The built-in differentiation features of Cheetah, allow Cheetah models to be used as **physics-based priors for Bayesian optimisation**.
 - Cheetah's implementation based on PyTorch allows it to seamlessly integrate with **modular neural network surrogate models**. This allows, for example, for **fast computation of collective effects**. Moreover, this integration preserves gradients.

Start using Cheetah yourself!

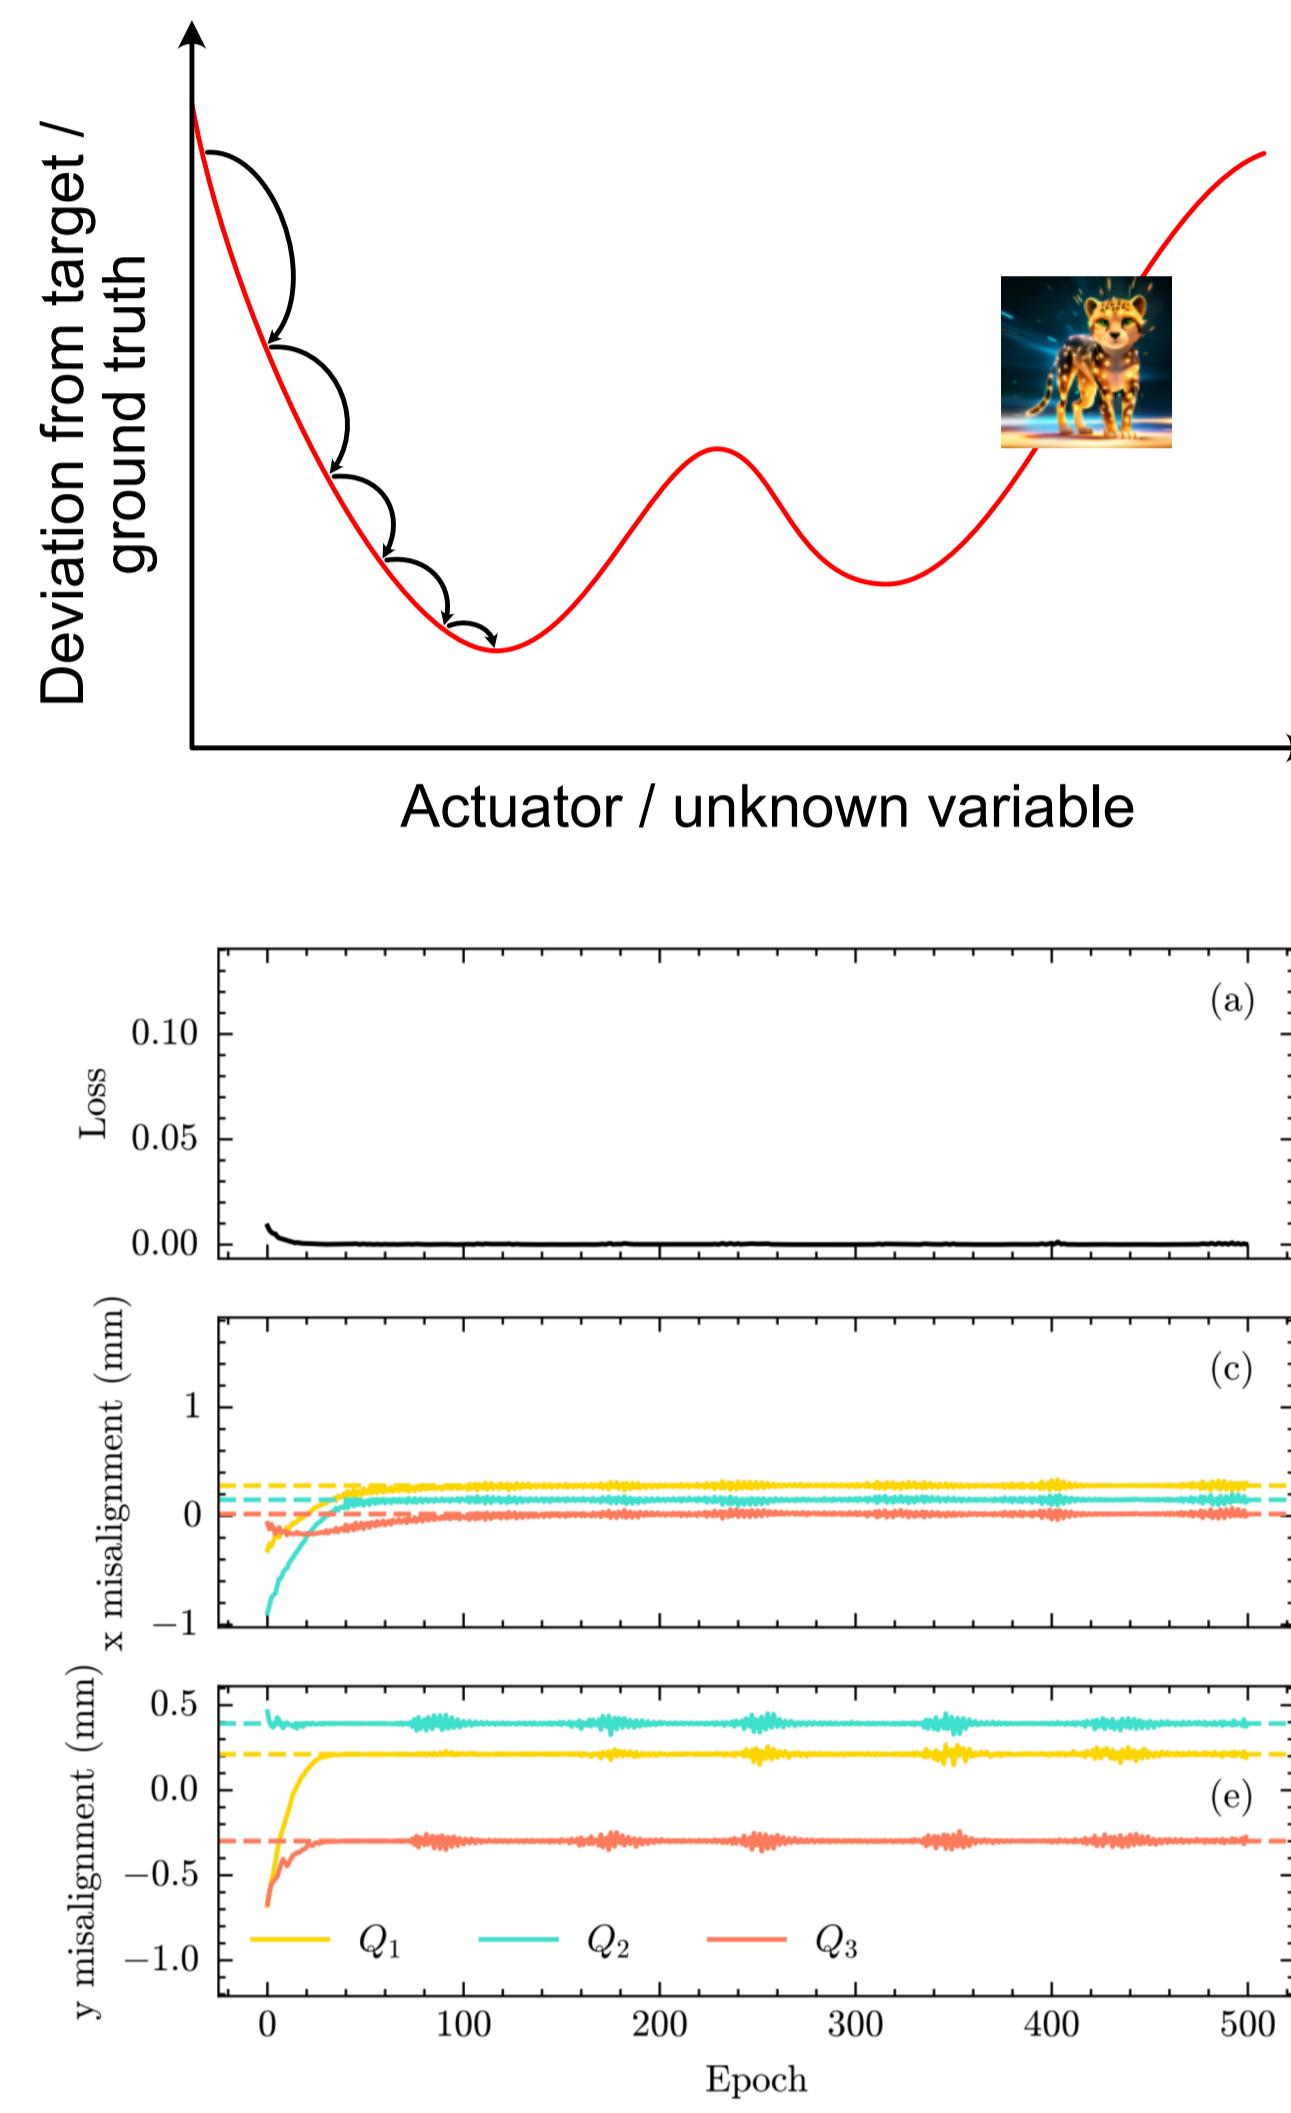
- Cheetah is available on [GitHub](https://github.com/desy-ml/cheetah) (<https://github.com/desy-ml/cheetah>) and can be installed through `pip install cheetah-accelerator`.

*Corresponding author: jan.kaiser@desy.de

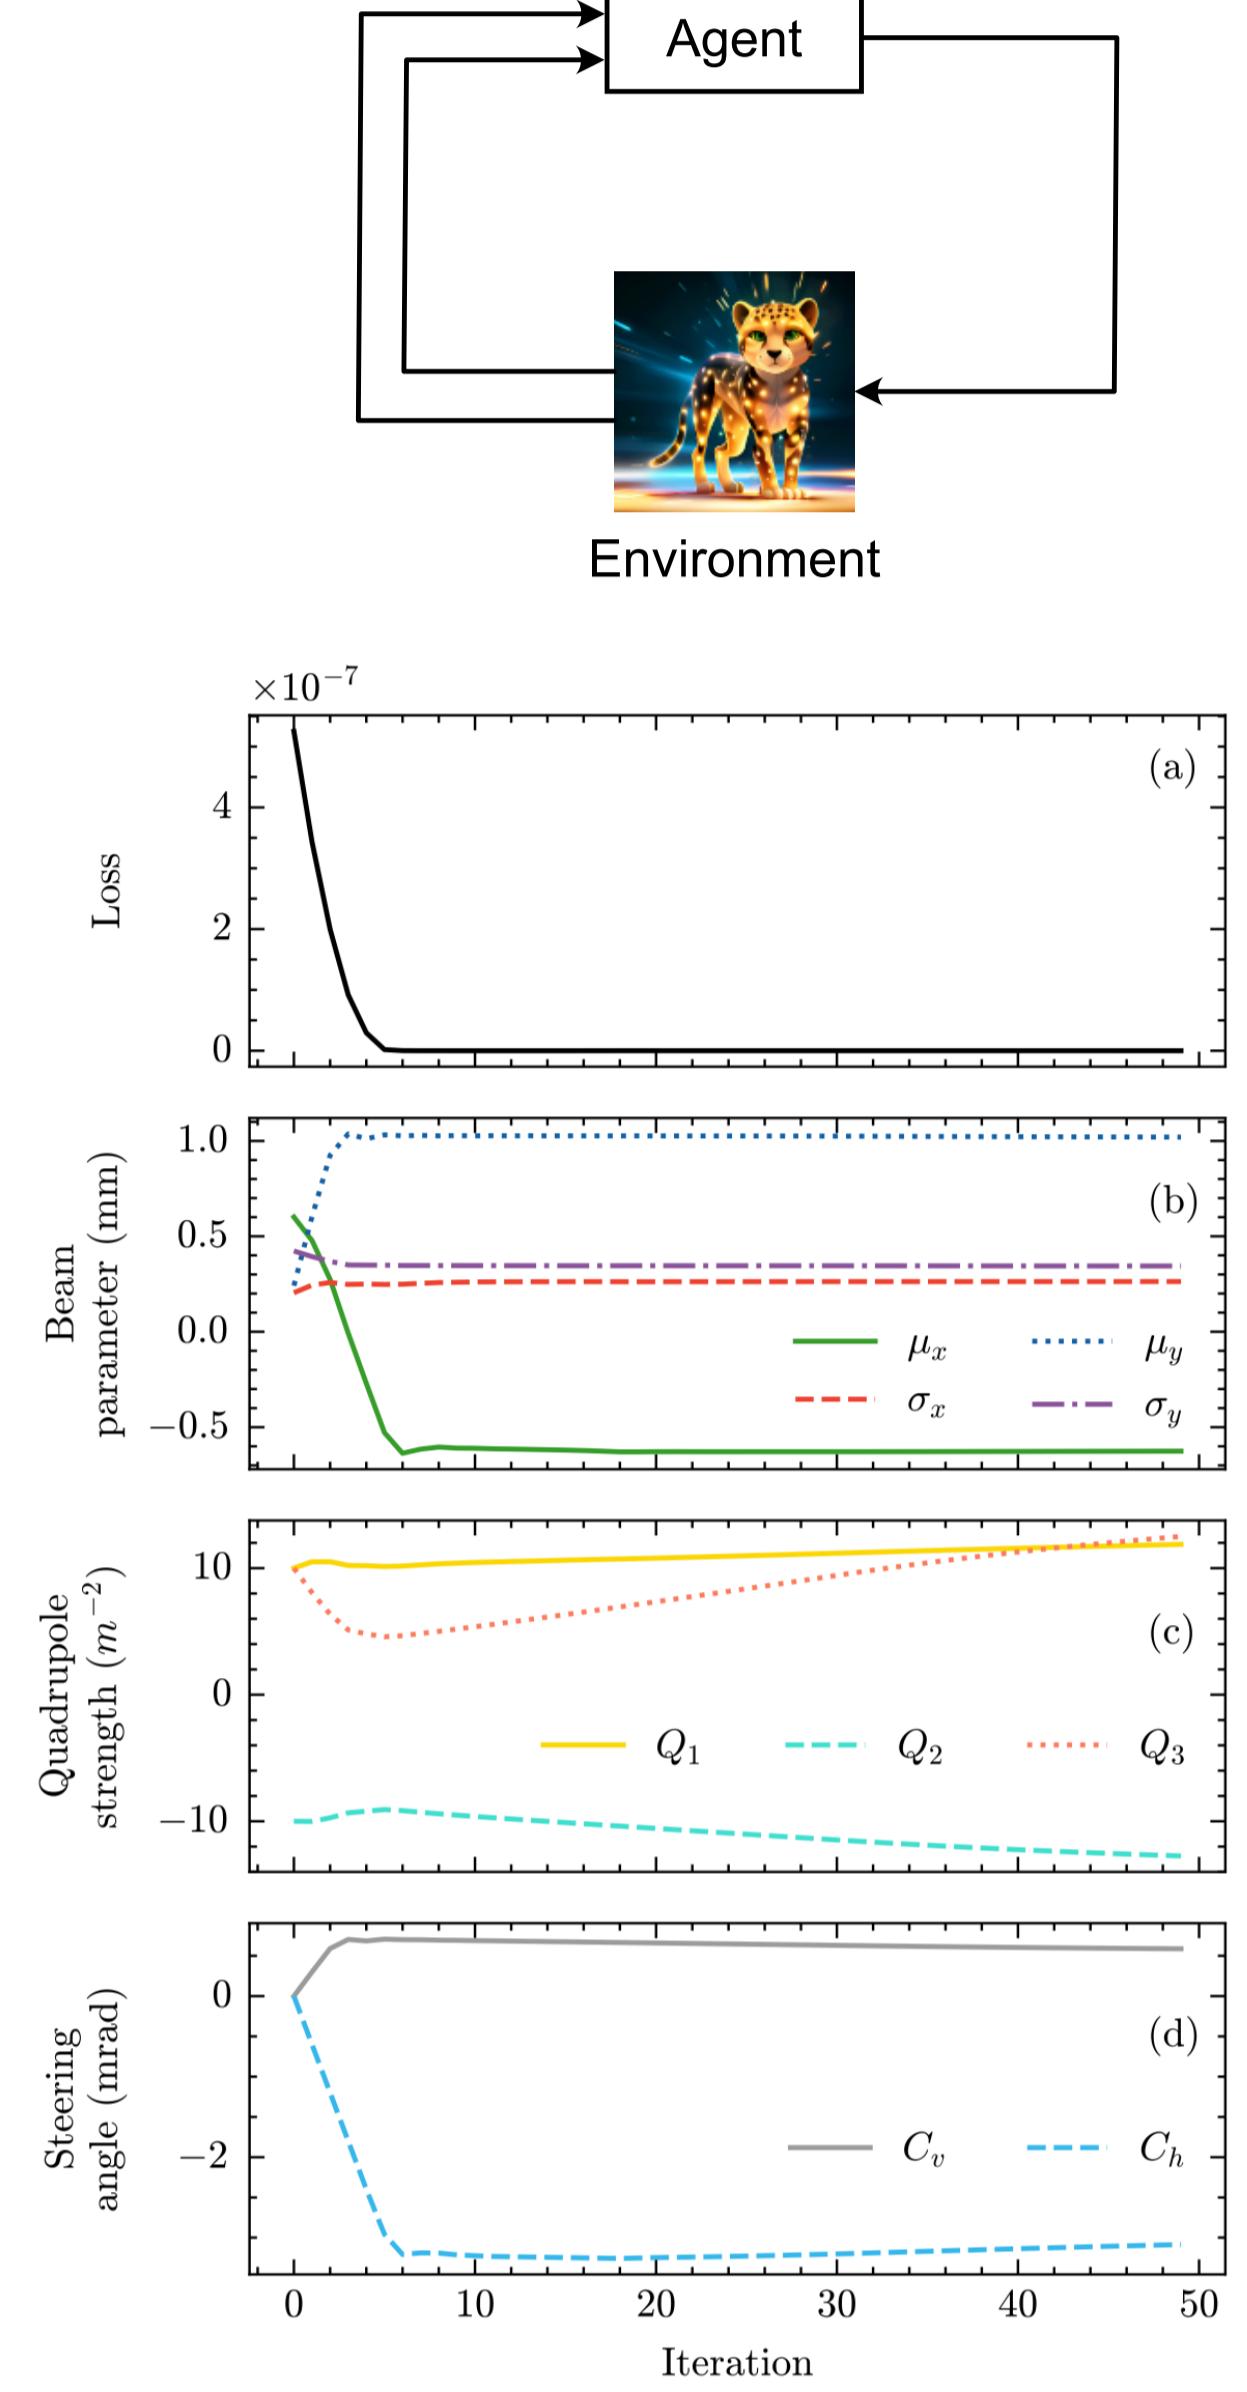


Scan me to download
the full paper!

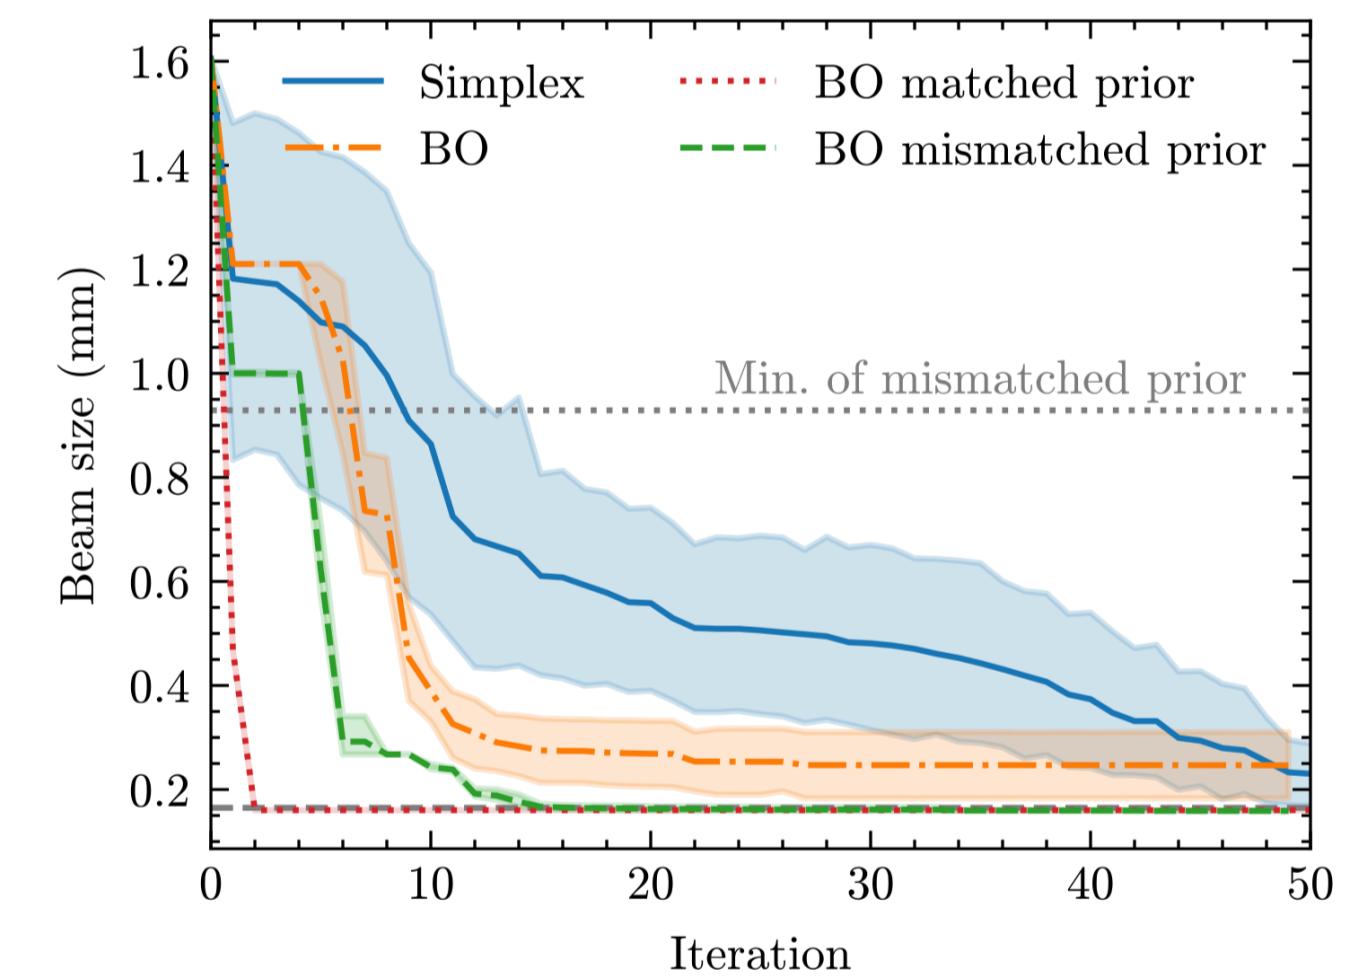
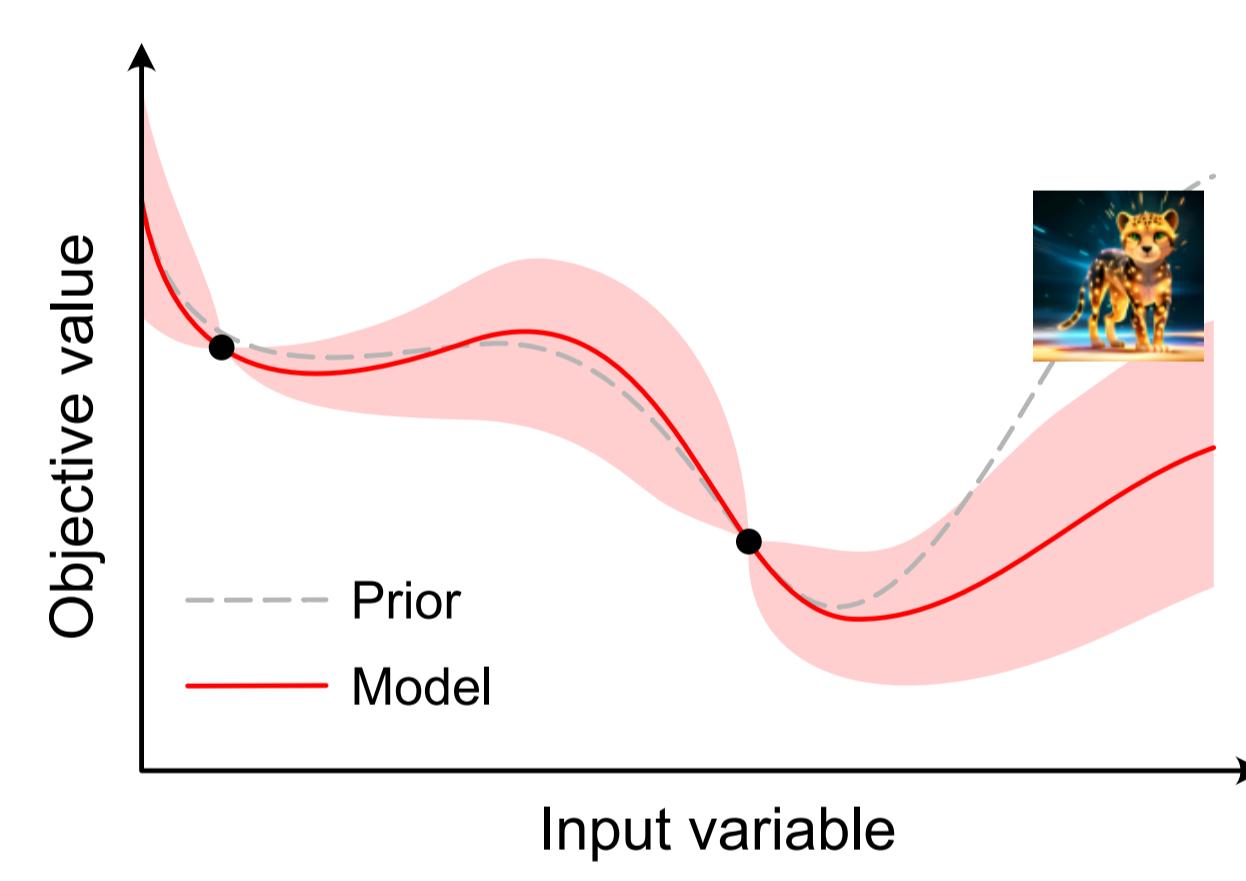
Gradient-based diagnostics



RLO tuning example



BO prior example



Neural network space charge surrogate example

